

# OOV Handling by Learning Subword using CNN based N-grams

Yonathan Purbo Santosa

Institut für Informatik  
Rheinische Friedrich-Wilhelms-Universität Bonn

November 13, 2019

# Overview

Introduction

Objective

CNN N-grams for OOV handling

Evaluation on Downstream Tasks

Conclusion and Future Works

# Introduction

## Why OOV handling is needed?

- ▶ Word embedding prefers large corpus [1]
- ▶ Language itself is creative and changing overtime [2, 3];
- ▶ Some word embedding has no OOV handling method [4, 5, 6]; and
- ▶ Improvements should be able to be achieved over using random or unknown embedding for OOV.

# Introduction

## Previous state-of-the-art Mimick

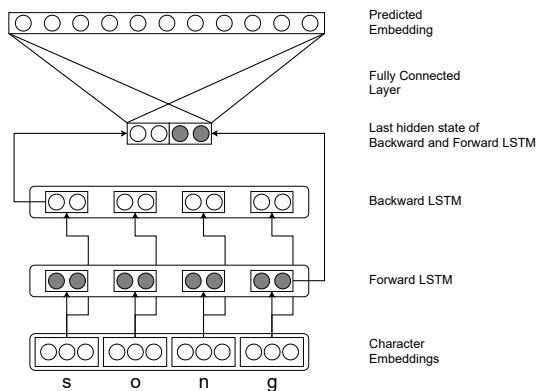


Figure: MIMICK Architecture

# Introduction

## What could be the problem with Mimick?

- ▶ MIMICK used the last hidden state of bi-LSTM to infer OOV embedding [7];
- ▶ By default, LSTM has a cell gate that could split the graph of the hidden unit when the input or previous state triggers it
- ▶ The last hidden state only encoded the recent sequence
- ▶ There is evidence that CNN can performs better than RNN and LSTM in sequence modeling [8]

# Objective

- ▶ Similar with MIMICK, this problem is tackled from quasi-generative perspective
- ▶ Instead of using bi-LSTM, CNN that based on n-grams can be used to learn character sequence [9].
- ▶ CNN N-grams can process different length sequence using max-over-time pooling.
- ▶ For each kernel  $K$ , an input that has similar features to the kernel can still gives a high response.

# CNN N-grams for OOV handling

- ▶ Given a word and its embedding  $(w_i, e_i)$
- ▶ The word is transformed into sequence of character embedding  $g_i \in \mathcal{G}$

$$h : w_i \rightarrow [c_1, c_2, \dots, c_n] \rightarrow [g_1, g_2, \dots, g_n]$$

- ▶ The sequence of character embedding then processed by the OOV handling model to predict the embedding  $\tilde{e}_i$

$$OOV_{model}(w_i) = \tilde{e}_i$$

# CNN N-grams for OOV handling




Figure: CNN N-grams for OOV Handling Architecture

- ▶ [2, 3, 4, 5, 6, 7]-grams were used as the feature extractor
- ▶ ReLU activation function was used before the max-over-time pooling



# CNN N-grams for OOV handling

- ▶ *Hardtanh* is used for predicting embedding
- ▶ The predicted embedding was used for calculating the deviation from the original embedding  $e_i$  by using MSE

$$\mathcal{L} = \frac{1}{2} \|e_i - \tilde{e}_i\|_2^2$$

Both model trained with three word embedding:

- ▶ Polyglot ( $\sim 100k$  words)
- ▶ Word2vec first 40k with reduction ( $\sim 35k$  words). If the word contain underscore “\_” or “http” then it will be removed
- ▶ Dict2vec 100 dimension ( $\sim 50k$  words)

# CNN N-grams for OOV handling

Table: OOV Handling Model Parameters

Hyperparameter	Mimick	CNN
Train-Val split	80%; 20%	
Batch size	64	
Epoch	100	
Momentum	0.5	
Learning Rate ( $\eta$ )	[0.01; 0.1]	0.1
Dropout	0	0.5
Num features	[50; 100; 200; 300]	[20; 50; 100; 150]

# Evaluation on Downstream Tasks

## Part-of-Speech Tagging

- ▶ Using brown POStagging datasets from NLTK<sup>1</sup>
- ▶ Following POS-tagger model from Ling et al. (2015) [10]
- ▶ Sentence length limited to 5 words
- ▶ There are 3 different setup for POS-tagger: Freeze, Continuous, Continuous with Trainable Embedding
- ▶ LogSoftmax is used for classifying the tags

---

<sup>1</sup>Available at <https://www.nltk.org/>

# Evaluation on Downstream Tasks

## Part-of-Speech Tagging

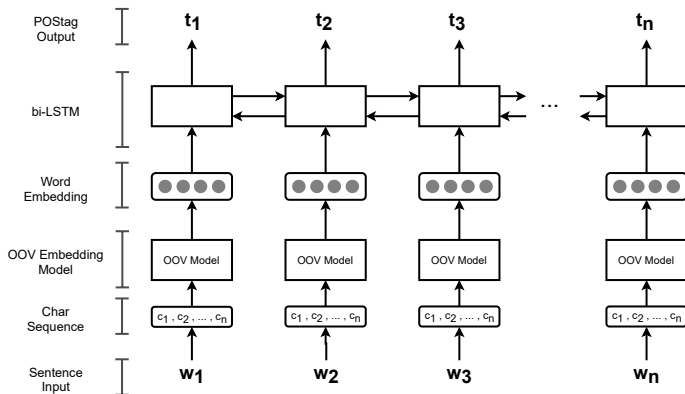


Figure: POS-tagger Model

# Evaluation on Downstream Tasks

## Part-of-Speech Tagging Result Random


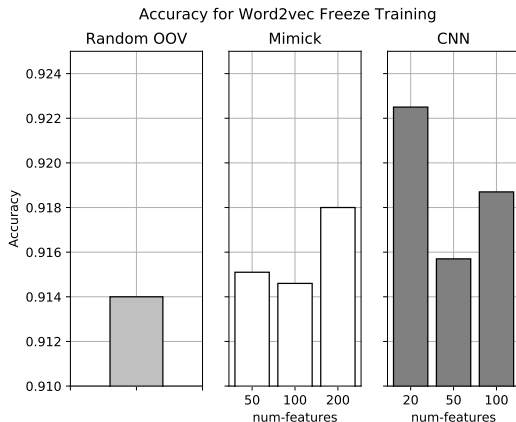


Figure: Random OOV Embedding


# Evaluation on Downstream Tasks

## Part-of-Speech Tagging Result Freeze




# Evaluation on Downstream Tasks

## Part-of-Speech Tagging Result Freeze



# Evaluation on Downstream Tasks

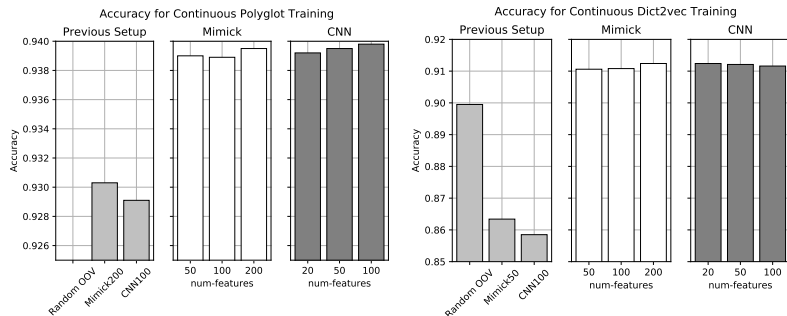
## Part-of-Speech Tagging Result Continuous





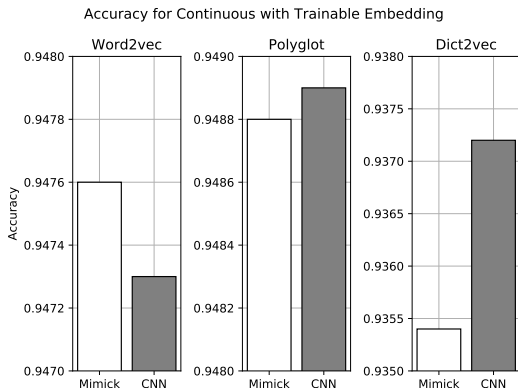
# Evaluation on Downstream Tasks

## Part-of-Speech Tagging Result Continuous



# Evaluation on Downstream Tasks

## Part-of-Speech Tagging Result Continuous with Trainable Embedding



# Evaluation on Downstream Tasks

## Word Similarity

- ▶ Using 12 different datasets to increase the data test size
- ▶ Only OOV embedding was inferred from the OOV handling model
- ▶ The cosine distance of a given pairs was calculated and compared with the dataset with Spearman's rank correlation coefficient.

# Evaluation on Downstream Tasks

## Word Similarity Results

Table: Word Similarity Task Results (word2vec)

Dataset	OOV	Invocab	OOV Ratio	Mimick*	CNN*
card	989	317	75.73%	<b>52.27</b>	27.25
mc30	4	35	10.26%	748.11	<b>758.79</b>
men	83	668	11.05%	<b>672.54</b>	664.64
mturk287	97	402	19.44%	518.15	<b>519.76</b>
mturk771	18	1095	1.62%	<b>657.93</b>	657.04
rg65	2	46	4.17%	709.77	<b>732.61</b>
rwstanford	1494	1457	50.63%	<b>233.83</b>	217.47
simlex	20	1008	1.95%	422.77	<b>425.90</b>
simverb	90	737	10.88%	<b>302.37</b>	292.91
verb143	4	113	3.42%	467.60	<b>478.25</b>
wordsim	15	422	3.43%	<b>658.92</b>	648.25
yp130	6	141	4.08%	<b>460.32</b>	443.22
average				<b>492.05</b>	488.84

\* multiplied by 1000

# Evaluation on Downstream Tasks

## Word Similarity Results

Table: Word Similarity Task Results (polyglot)

Dataset	OOV	Invocab	OOV Ratio	Mimick*	CNN*
card	864	442	66.16%	<b>128.93</b>	114.11
mc30	1	38	2.56%	605.25	605.25
men	14	737	1.86%	490.57	<b>492.17</b>
mturk287	76	423	15.23%	443.31	<b>458.81</b>
mturk771	3	1110	0.27%	<b>432.48</b>	432.20
rg65	1	47	2.08%	531.59	<b>524.77</b>
rwstanford	999	1952	33.85%	272.33	<b>290.78</b>
simlex	4	1024	0.39%	232.20	<b>234.16</b>
simverb	53	774	6.41%	<b>137.01</b>	134.42
verb143	0	117	0.00%	335.81	335.81
wordsim	0	437	0.00%	412.83	412.83
yp130	5	142	3.40%	<b>44.76</b>	44.62
average				338.92	<b>339.99</b>

\* multiplied by 1000

# Evaluation on Downstream Tasks

## Word Similarity Results

Table: Word Similarity Task Results (Dict2vec)

Dataset	OOV	Invocab	OOV Ratio	Random*	Mimick*	CNN*
card	828	478	63.40%	48.07	80.89	<b>95.32</b>
mc30	0	39	0.00%	847.57	847.57	847.57
men	1	750	0.13%	713.16	723.63	<b>723.89</b>
mturk287	2	497	0.40%	652.27	<b>655.32</b>	653.13
mturk771	0	1113	0.00%	683.91	683.91	683.91
rg65	0	48	0.00%	832.86	832.86	832.86
rwstanford	619	2332	20.98%	214.60	<b>403.79</b>	400.27
simlex	3	1025	0.29%	454.80	<b>460.66</b>	459.87
simverb	24	803	2.90%	375.15	390.09	<b>393.39</b>
verb143	0	117	0.00%	187.82	187.82	187.82
wordsim	18	419	4.12%	642.71	718.72	<b>723.72</b>
yp130	2	145	1.36%	577.76	621.38	<b>621.75</b>
average				519.22	550.55	<b>551.96</b>

\* multiplied by 1000

# Conclusion

- ▶ Machine learning can be used to generate OOV embedding
- ▶ Compared to random OOV embedding, machine learning method has higher performance in downstream tasks
- ▶ CNN generally performs better than bi-LSTM for POS-tagging and word similarity tasks
  - ▶ MIMICK has higher accuracy in POS-tagging: when the pre-trained embedding & the OOV handling model are non-trainable (*out-of-the-box*)
  - ▶ CNN N-grams has higher accuracy in POS-tagging: when both pre-trained embedding and the OOV handling model are trainable
  - ▶ Despite that MIMICK performs better *out-of-the-box* for POS-tagging, this was not the case for word similarity task

## Future Works

- ▶ A function that can generate an entire embedding
- ▶ Save space (e.g. word2vec 3 million words and phrases is around 2GB)
- ▶ Save computational needs for new entries cases



Thank you for your attention!

# References



A. Kutuzov and M. Kunilovskaya (2018)

Size vs. structure in training corpora for word embedding models: Araneum russicum maximum and russian national corpus  
*Analysis of Images, Social Networks and Texts 2018*  
Springer International Publishing



Julie C. Forrester (2008)

A Brief Overview of English as a Language in Change  
*Canadian Center of Science and Education 2008*, Vol. 4(4)



Daniel Jurafsky and James H. Martin (2009)

Speech and Language Processing (3rd Edition)  
*Prentice-Hall, Inc.*



Rami Al-Rfou, Bryan Perozzi, and Steven Skiena (2010)

Polyglot: Distributed Word Representations for Multilingual NLP  
*Proceedings of the Seventeenth Conference on Computational Natural Language Learning*  
Association for Computational Linguistics



Julien Tissier, Christophe Gravier, and Amaury Habrard (2017)

Dict2vec : Learning Word Embeddings using Lexical Dictionaries  
*Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, 254–263  
Association for Computational Linguistics

# References



Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean (2013)  
Efficient Estimation of Word Representations in Vector Space  
*1st International Conference on Learning Representations, ICLR 2013*



Yuval Pinter, Robert Guthrie, and Jacob Eisenstein (2017)  
Mimicking Word Embeddings using Subword RNNs  
*Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*



Shaojie Bai and J. Zico Kolter and Vladlen Koltun (2018)  
An Empirical Evaluation of Generic Convolutional and Recurrent Networks for Sequence Modeling  
*Computing Research Repository (CoRR)*



Yoon Kim (2014)  
Convolutional Neural Networks for Sentence Classification  
*Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing*, 1746–1751



Wang Ling, Chris Dyer, Alan W. Black, Isabel Trancoso, Ramon Fernandez, Silvio Amir, Luis Marujo, and Tiago Luis (2015)  
Finding Function in Form: Compositional Character Models for Open Vocabulary Word Representation  
*Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing* 1520–1530